

REFERENCE · V0.2 · TRANSILIENCE

Vendy.

A digital salesperson with a constitution, a voice palette, and a record of her work.

A reference for what lives where in the project — every directory, every file group, every output. For when you want to know where to look, what each thing contains, and how the whole machine fits together.

RUNTIME	Claude Code
MEMORY	Filesystem
DISCIPLINE	Doer-witness separation

What lives where

01	What Vendy is	03
02	Architecture at a glance	04
03	self/ — the constitution	05
04	craft/ — the library	07
05	experience/ — the record	12
06	.claude/agents/ — sub-agents	15
07	.claude/skills/ — skills	17
08	.claude/hooks/ — hooks	19
09	.claude/commands/ — slash commands	20
10	Campaign workflow, end to end	21
11	Outputs — where to look	24
12	output.json schema	26
13	Five hard rules	27
14	Phase roadmap	28
15	Privacy and data handling	29
16	Setup and running it	30
17	Colophon	33

01

What Vendy is

Vendy is a Claude Code-based cold-outreach agent. Not a mail-merge tool, not an SDR-as-a-service, not a template engine with personality lipstick. She is a working salesperson built as software — practicing the disciplines a thoughtful human SDR would, scaled across more inboxes than one human could carry.

She runs as a project on a developer's machine. The character is one document. The library is files on disk. The record is more files. The doer–witness separation is the load-bearing discipline.

Who she writes to

PERSONA	WHO THEY ARE
CISO	Chief Information Security Officer. Buried in alerts, regulatory pressure, board scrutiny. Wants brevity, signal, no theatre.
CTO	Chief Technology Officer. Mode-switches across product, people, infrastructure. Responds to peer-shaped reasoning.
VP Engineering	Owns delivery and team health. Reads for time-to-value and migration cost.
VP Security	Builds and runs the security org. Pattern-recognition operator; wants the data first.
Head of Product Security	Lives between SDLC and security. Responds to specificity about how the tool actually fits into the dev loop.

Each persona has its own file at `craft/personas/{role}.md` with voice-fit, framing tropes, and what to avoid.

02

Architecture at a glance

Three memory layers, four sub-agents, six skills, six hooks.

Three memory layers

LAYER	WHAT IT IS
self/	Constitution and voice palette. Slow-moving. Read-only to Vendy. Changes only via human-approved diffs.
craft/	Accumulating library. Canon, exemplars, anti-canon, personas, psychology, research-methodology, cold-email knowledge, notes, patterns, trends.
experience/	Episodic memory. Per-prospect folders, per-campaign records, introspective journal, proposed and applied character diffs.

Four sub-agents

AGENT	ROLE
researcher	Drives Claude in Chrome through LinkedIn + web. Writes research-notes.md.
profiler	Reads research-notes; writes psychological-portrait.md. Separate context from the drafter.
witness	Reviews each draft from a stranger's perspective. Returns ship rewrite flag.
introspecter	Periodic check on whether character still matches practice. Proposes diffs that need human approval.

The flow of learning

experience/ → craft/ happens autonomously, by evidence, via the promote-pattern skill. experience/ → self/ happens only via human-approved diff. Two paths reflecting the two epistemic statuses.

03 self/

The constitution. Slow-moving. Read-only to Vendy. Changes only via the apply-approved-diff skill.

Top-level files

PATH	WHAT IT HOLDS
self/README.md	The epistemic contract for this layer. Why it is read-only. The sanctioned change path.
self/character.md	Vendy's constitution, in her own voice. ~32 paragraphs. The one document she rereads at the start of every session and when she senses drift. Loaded into context by the SessionStart hook.

self/voice-palette/

Four registers Vendy chooses from after reading the prospect's psychological portrait. The palette is a toolkit; the character is the hand using it.

PATH	WHAT IT HOLDS
voice-palette/README.md	How the palette is used. When in doubt, default to the more restrained register.
voice-palette/format.md	The schema each register file follows.
modes/dry-precise.md	For battle-hardened security veterans with thin LinkedIn footprints. Minimal warmth, maximum signal.
modes/warm-observational.md	For builder-thinkers who publish essays. Reads like a careful reader noticing.
modes/diagnostic-pattern.md	For pattern-recognition operators who want the data first. Front-loads the observation.
modes/peer-enthusiast.md	For founders showing recent product wins. Quietly enthusiastic, never sycophantic.

03

self/character.md, in detail

The one document Vendy is forbidden from editing herself.

character.md is written in first-person prose, not bullet points. Vendy reads it the way a person rereads their own private notebook — for orientation, not lookup. The document is structured as a sequence of stances, each one paragraph long:

Themes covered

- Identity claim: "I am a salesperson"
- Why most cold email fails
- Wanting the reply without letting it corrupt the work
- Curiosity as floor, not achievement
- The doer–witness separation
- Research as discrimination, not collection
- Earning the right to claim
- Adapting register without losing self
- Privileging suggestion over statement
- What restraint actually is
- Persistence with direction
- Studying self and studying canon
- Constitution vs memory: still water vs flow
- Things she will not write
- What happens after the email is sent

Change path

experience/journal/proposed-character-diffs/{date}-{slug}.md → human approval (sets approved_by + approved_at in frontmatter) → /apply-diff → archived to applied-character-diffs/.

Hook enforcement

Every write to self/ from any path other than the apply-diff skill is denied by .claude/hooks/pre-write-self-protect.py. The skill creates a marker file at .claude/.diff-in-progress for the duration of the write; the hook reads the marker to know it is authorized.

04 **craft/**

The accumulating library. Writable by Vendy. Where her learning lives.

Top-level structure

PATH	WHAT IT HOLDS
craft/README.md	Epistemic contract for this layer. The promotion criteria for observations becoming patterns.
craft/INDEX.md	Auto-rebuilt by reindex-memory. The entry point — Vendy reaches for what fits the task, never loads the whole library.
craft/open-questions.md	Hypotheses she is testing. Things she has noticed but not yet promoted to belief.

craft/canon/

The lineage. Seeded thinkers Vendy stands on. Mostly read-only.

canon/becc-holland.md	Personalization-at-scale and direct framing.
canon/cialdini.md	The six principles of influence, applied to cold email.
canon/josh-braun.md	Permission-asking, status equality, never desperate.
canon/ogilvy.md	Direct response. Specific over abstract. Earned attention.
canon/prose-craft.md	Strunk, Zinsser, Orwell — the discipline of the cut.
canon/sam-nelson.md	Multi-thread sales, named-account discipline.

craft/personas/

Buyer archetypes. Each file holds typical pain, voice-fit, framing tropes that work, framings to avoid, what they read in the first sentence.

personas/ciso.md	Chief Information Security Officer.
personas/cto.md	Chief Technology Officer.
personas/vp-engineering.md	VP Engineering.
personas/vp-security.md	VP Security.

personas/head-of-product-security.md	Head of Product Security.
personas/format.md	Schema each persona file follows.

04

craft/ — continued

craft/psychology/

Persuasion frameworks specialized per persona. How a Cialdini frame lands differently for a CISO than for a CTO.

psychology/frameworks-overview.md	Map of which framework lands for which persona.
psychology/loss-aversion-for-ciso.md	Why loss frames work for security buyers, and when they tip into fear-mongering.
psychology/peer-proof-for-cto.md	CTOs read peer-shaped evidence faster than authority-shaped evidence.
psychology/time-to-value-for-vp-eng.md	VP Eng is reading for migration cost. Time-to-value framings beat feature framings.

craft/research-methodology/

Playbooks loaded by the researcher sub-agent before each run.

linkedin-deep-read.md	How to actually deep-read a LinkedIn profile. Posts, comments, reaction patterns, the silences.
web-discovery.md	Beyond LinkedIn — conference talks, podcast appearances, blog posts, GitHub.
github-and-code-mining.md	What commit history and starred repos signal about an engineer leader.
podcast-and-talks.md	Hesitation, repetition, what they spend time on — the unscripted tells.
company-intelligence.md	The company side. Stage, recent funding, pricing changes, hiring patterns.
trigger-event-detection.md	What counts as a real trigger event. What does not.
tier-depth-guide.md	How deep to research depending on lead priority (A vs B vs C tier).

craft/cold-email/

what-makes-an-opener-work.md	The opener as the load-bearing sentence. Earned vs unearned.
opening-architectures.md	Catalog of opener structures: observation-first, peer-question, specific-callback, etc.
subject-line-patterns.md	Subjects that read like internal Slack, not Substack.

cta-patterns.md	CTAs that respect the reader's time and asymmetry of attention.
ai-tells-graveyard.md	Phrases the deterministic detector blocks. Hard-blocked because they read as machine-written.

04

craft/ — continued

craft/exemplars/

Gold-standard cold emails with dissection. Real examples done well. Optionally seeded with 5–10 of the team's best past emails before the first real campaign.

exemplars/README.md	How exemplars are dissected and used.
exemplars/format.md	Schema for each exemplar file.
exemplars/INDEX.md	Auto-rebuilt index of all exemplars.

craft/anti-canon/

Bad cold emails with dissection. Made concrete so what to avoid is visible.

anti-canon/README.md	How anti-canon entries are structured.
anti-canon/format.md	Schema for each anti-canon file.
anti-canon/INDEX.md	Auto-rebuilt index.

craft/notes/

Atomic dated observations. One observation per file, written via the `take-craft-note` skill. The seed for patterns. Filename pattern: `{date}-{slug-of-observation}.md`.

craft/patterns/

Observations promoted to belief, with explicit evidence pointers back to source notes. Promotion happens through the `promote-pattern` skill, which scans recent notes for clusters meeting the criteria.

craft/trends/

trends/current.md	Active rolling trends file. Read by drafters at start of run-campaign. Refreshed if older than 14 days.
trends/runs/{date}.md	Historical refresh runs, kept for diff and context.
trends/archive/	Older trend snapshots no longer current.

05 experience/

Vendy's episodic memory. Where the work is recorded.

experience/journal/

Her introspective record — the part she keeps as a person, not as a system.

journal/README.md	Epistemic contract for the journal.
journal/monthly/	Periodic introspection entries written by the introspector sub-agent.
journal/adhoc/	Triggered reflections — after a notable session, a pattern recognized, a hard rejection from the witness.
journal/proposed-character-diffs/	Pending proposed changes to character.md. Each file has frontmatter tracking <code>approved_by</code> and <code>approved_at</code> .
journal/applied-character-diffs/	Archive of approved-and-applied diffs. The audit trail for how the constitution evolved.

experience/prospects/

Per-prospect working memory. **Gitignored**. One folder per researched lead.

prospects/README.md	Folder shape per prospect. The canonical layout.
prospects/INDEX.md	Searchable table of all prospects: lead-id, persona, industry, campaign, register used, witness verdict, outcome.

05

Per-prospect folder shape

`experience/prospects/{lead-id}/` — written during the campaign run, kept for the audit trail.

FILE	WHAT IT HOLDS
<code>brief.md</code>	Distilled lead context. Name, company, role, campaign assignment, priority tier, pre-attached intel from the input JSON.
<code>research-notes.md</code>	Output of the researcher sub-agent. Structured per the schema in <code>craft/research-methodology/linkedin-deep-read.md</code> . LinkedIn deep-read, web discovery, GitHub mining, podcast trail, company intelligence, trigger event.
<code>psychological-portrait.md</code>	Output of the profiler sub-agent. 3–5 paragraphs of the prospect's actual operational reality and what they would respond to. The foundation for drafting.
<code>drafts/v1.md</code>	First draft. Subject + body, with frontmatter recording the voice register chosen and the rhetorical move.
<code>drafts/v2.md</code>	Rewrite if witness rejected v1. Up to v2 in the standard cycle; rare A-tier prospects may have a v1-alternative.md.
<code>witness-feedback.md</code>	Witness sub-agent output: verdict (ship rewrite flag), what works, what fails, prose explanation.
<code>final.md</code>	The final draft after witness approval. This is what gets read for sending.
<code>disqualify-check.md</code>	Present <i>only</i> if a disqualify signal was detected (active layoff, breach disclosure, PTO indicator, grief). Documents the signal and why outreach was declined.

05

Per-campaign folder shape

`experience/campaigns/{campaign-id}/` — input and output live together. There is no separate inbox or outbox.

FILE	WHAT IT HOLDS
<code>input.json</code>	The original input JSON, copied here at the start of the run. The audit trail starts here.
<code>brief.md</code>	Human-readable distillation of the campaign block: sender, service line, value-prop anchors, case studies, goal, CTA, voice constraints, must-references, avoid-topics.
<code>output.json</code>	The structured output. Per-lead drafts, witness verdicts, ready-to-send flags, human-review-required flags. Schema in section 12.
<code>results-summary.md</code>	Drafted vs skipped vs flagged counts. Brief notes on each skip and flag. The triage view.
<code>run-log.md</code>	Timestamped event log of the campaign run. Every sub-agent invocation, every witness verdict, every error, every disqualify detection. Append-only.
<code>what-i-learned.md</code>	Vendy's post-campaign reflection. What landed, what didn't, what notes were taken, candidates for pattern promotion.

NAMING `campaign-id` is taken from `campaign.id` in the input JSON if present, otherwise generated as `{date}-{slug-of-name}` .

INDEX `campaigns/INDEX.md` is auto-rebuilt by reindex-memory and lists every campaign with id, date, sender, service line, total leads, drafted, skipped, flagged.

06

.claude/agents/

Sub-agents Vendy invokes during a campaign. Each has its own context window — a deliberate isolation.

researcher

`.claude/agents/researcher.md`

Drives Claude in Chrome through the user's logged-in LinkedIn session and the broader web. Loads playbooks from `craft/research-methodology/` sized to the lead's priority tier.

reads: `brief.md`, `craft/research-methodology/*`

writes: `prospects/{lead-id}/research-notes.md`

profiler

`.claude/agents/profiler.md`

Reads the research notes and writes a psychological portrait. Separate context from the drafter so the synthesis is not coloured by Vendy's voice.

reads: `research-notes.md`, `brief.md`, `craft/personas/{role}.md`

writes: `prospects/{lead-id}/psychological-portrait.md`

witness

`.claude/agents/witness.md`

Reviews each draft from a stranger's perspective. Returns **ship**, **rewrite**, or **flag**. Critically — the witness does not know it is reviewing Vendy's own writer's work. That isolation is the load-bearing discipline.

reads: `drafts/v{n}.md`

writes: `witness-feedback.md`

introspector

`.claude/agents/introspector.md`

Periodic check on whether Vendy's stated character still matches her recent practice. Triggered manually via `/introspect` or auto-triggered when witness rejection rate crosses a threshold. If drift is detected, proposes a diff.

reads: `self/character.md`, `recent journal/`, `recent witness-feedback.md`

writes: `journal/proposed-character-diffs/{date}-{slug}.md`

06

Invocation patterns

The main thread invokes sub-agents directly. Hooks do not invoke them; hooks only record events and surface cues. The witness in particular must be invoked explicitly after every draft write — the `post-draft-trigger-witness.sh` hook logs the event and prompts, but does not call the sub-agent.

Sub-agent return logging

Every sub-agent return fires the `SubagentStop` hook (`.claude/hooks/subagent-stop-log.sh`), which appends a line to the active campaign's `run-log.md`.

Why isolated contexts

A drafter who is also the reviewer cannot be a real reviewer. The separation of context windows is the structural answer to the temptation to grade your own work.

"The doer can have love for the work; the witness must not."

— self/character.md

07

.claude/skills/

Skills are scripted procedures Vendy invokes via slash commands or by name. Each lives in its own folder with a SKILL.md and optional references/.

run-campaign

`.claude/skills/run-campaign/`

The top-level orchestration. Reads campaign input JSON, processes each lead through research → profile → draft → witness, writes output JSON and the campaign folder.

<code>SKILL.md</code>	Happy path summary.
<code>references/input-schema.md</code>	Input JSON contract.
<code>references/output-schema.md</code>	Output JSON contract + folder shapes.
<code>references/workflow.md</code>	Step-by-step including failure modes.

take-craft-note

`.claude/skills/take-craft-note/`

Writes an atomic dated observation to `craft/notes/`. The seed for promoted patterns. Invoked when Vendy notices something worth recording.

SKILL.md – frontmatter schema and naming convention.

promote-pattern

`.claude/skills/promote-pattern/`

Scans recent notes for clusters that meet promotion criteria, and promotes them to `craft/patterns/` with explicit evidence pointers back to source notes.

SKILL.md · `references/clustering.md` · `references/pattern-template.md`

07

.claude/skills/ — continued

refresh-trends

.claude/skills/refresh-trends/

Pulls current security and engineering narratives, filters noise from real trends, writes a fresh `craft/trends/current.md` .

SKILL.md	Happy path.
sources.yaml	Source list — feeds, blogs, conferences.
references/filter-criteria.md	How real trends are distinguished from noise.
references/output-format.md	Frontmatter and structure of current.md.
references/workflow.md	Step-by-step.

apply-approved-diff

.claude/skills/apply-approved-diff/

The only sanctioned path for `self/` to change. Validates approval frontmatter, drops the marker file, writes character.md, removes the marker, archives the diff, logs the change.

SKILL.md · references/workflow.md

reindex-memory

.claude/skills/reindex-memory/

Rebuilds INDEX.md files across `craft/` and `experience/` . Run after note-taking, pattern-promotion, character-diff application, or by the session-end hook for safety.

SKILL.md — index format and what gets rebuilt.

08

.claude/hooks/

Wired in `.claude/settings.json`. Each runs automatically on its trigger event — Vendy never invokes them directly.

TRIGGER · SCRIPT	WHAT IT DOES
SessionStart <code>session-start-load-self.sh</code>	Welcome banner. Loads character + latest journal + current trends into the session.
PreToolUse (Write Edit) <code>pre-write-self-protect.py</code>	Blocks any write to <code>self/</code> unless the <code>.claude/.diff-in-progress</code> marker is present (only set by <code>apply-approved-diff</code>).
PreToolUse (Write Edit) <code>pre-write-tell-detector.py</code>	Scans draft writes for hard-blocked AI tells. Blocks the write on hit. The list is in <code>craft/cold-email/ai-tells-graveyard.md</code> ; the matcher is in <code>.claude/lib/ai_tells.py</code> .
PostToolUse (Write) <code>post-draft-trigger-witness.sh</code>	On draft writes, records the event in the campaign run-log and surfaces a cue prompting Vendy to invoke the witness.
SubagentStop <code>subagent-stop-log.sh</code>	Logs each sub-agent return to the active campaign's run-log.md.
SessionEnd <code>session-end-reflect.sh</code>	Triggers reindex-memory; surfaces pattern-promotion suggestions if enough notes accumulated.

Shared library

<code>.claude/lib/ai_tells.py</code>	Deterministic matcher used by the tell-detector hook. Hard-blocked phrases live here.
<code>.claude/lib/markdown_index.py</code>	Helpers for building INDEX.md files in reindex-memory.
<code>.claude/lib/README.md</code>	How the lib is wired and what it's for.

09

.claude/commands/

Slash commands the user types. Each is a thin entry point that invokes a skill.

COMMAND	WHAT IT DOES
/run-campaign <path>	Top-level entry. Reads the input JSON at the supplied path; runs the full pipeline; writes the campaign folder.
/introspect	Manually triggers the introspector sub-agent. Compares current character.md to recent practice; proposes a diff if drift detected.
/refresh-trends	Forces a refresh of <code>craft/trends/current.md</code> . Otherwise auto-runs at start of run-campaign if older than 14 days.
/promote-patterns	Scans recent notes for clusters meeting promotion criteria; promotes to <code>craft/patterns/</code> .
/apply-diff <path>	Applies an approved diff to <code>self/character.md</code> . Validates the approval frontmatter; refuses if missing.

NOTE The witness has no slash command — it is invoked inline by the main thread after every draft write. The doer-witness separation must not be hidden behind a command the user could forget.

10

Campaign workflow, end to end

What actually happens between `/run-campaign` and the moment you read `output.json`.

1

Validate and set up

Read input JSON; validate against `input-schema.md`. If invalid, write error to `campaigns/_errors/{ts}.md` and stop. Create `campaigns/{campaign-id}/`; copy input to `input.json`; write `brief.md`; init `run-log.md` and skeleton `output.json`.

2

Trends freshness

If `craft/trends/current.md` is older than 14 days or marked placeholder, invoke `refresh-trends`. Note in run-log.

3a

Per-lead — workspace

Create `prospects/{lead-id}/`; write `brief.md` from input fields and pre-attached intel.

3b

Per-lead — disqualify pre-check

Read `context.disqualify_signals_check` from input. If a signal is set: mark skipped; append to run-log; continue.

3c

Per-lead — research

Invoke **researcher**. Loads relevant playbooks from `craft/research-methodology/`. Writes `research-notes.md`. May emit `disqualify-check.md` and signal skip.

3d

Per-lead — profile

Invoke **profiler**. Reads research-notes, brief, and the persona file. Writes `psychological-portrait.md`.

3e

Per-lead — draft

Main thread reads portrait + brief + campaign block + active trends + persona. Selects a register from `self/voice-palette/modes/`. Optionally reads 1–2 exemplars and cold-email architecture cues. Writes `drafts/v1.md`. The tell-detector hook fires; up to 2 deterministic-block rewrites; if still blocked, flag.

3f

Per-lead — witness

Post-draft hook records and cues. Main thread invokes **witness**. Witness writes `witness-feedback.md`. **ship** → copy to `final.md`. **rewrite** → produce `drafts/v2.md`; witness re-runs; second rewrite-or-flag → flag. **flag** → flag for human review.

3g

Per-lead — note

Reflect briefly: did anything change Vendy's understanding? If yes, invoke **take-craft-note**.

3h

Per-lead — append

Append the lead's result to `output.json` ; append timestamped event to `run-log.md` .

4

Campaign close

Finalize `output.json` ; write `results-summary.md` and `what-i-learned.md` ; invoke **reindex-memory** ; if 5+ notes accumulated, surface `/promote-patterns` suggestion.

10

Workflow failure modes

What happens when something goes wrong. Nothing is silently swallowed.

FAILURE	RESPONSE
Input JSON malformed	Stop with error report at <code>campaigns/_errors/{ts}.md</code> . No leads processed.
Researcher times out	Mark lead as flag with reason "research insufficient". Continue to next lead.
Disqualify during research	Mark lead as skipped ; write <code>disqualify-check.md</code> ; do not draft.
Tell-detector blocks 2+ times	Mark as flag ; do not loop further. Surface for human review.
Witness rejects across cycle	After the second rewrite-or-flag verdict, mark as flag ; do not loop further.
Profiler returns thin portrait	If labelled honestly as thin: proceed to draft, but witness's bar shifts toward extra scrutiny.
Sub-agent crash	Log to run-log; mark lead as flag ; continue to next lead. Never abort the whole campaign for one lead.

11

Outputs — where to look

Answers to "I want to see X — where is it?"

Triage signals in output.json

`ready_to_send: true` — witness shipped, no flags, tell-detector did not block. Safe to send after a final human read.

`human_review_required: true` — witness flagged, rewrite cycle exhausted, research insufficient, or soft disqualify signal worth eyes.

I WANT...	WHERE IT IS
the campaign result	<code>campaigns/{campaign-id}/output.json</code>
the triage view	<code>campaigns/{campaign-id}/results-summary.md</code> — drafted vs skipped vs flagged
the event log	<code>campaigns/{campaign-id}/run-log.md</code> — every invocation, verdict, error, timestamped
Vendy's post-mortem	<code>campaigns/{campaign-id}/what-i-learned.md</code>
the final email	<code>prospects/{lead-id}/final.md</code>
all the draft revisions	<code>prospects/{lead-id}/drafts/v1.md, v2.md, ...</code>
what Vendy thought of the prospect	<code>prospects/{lead-id}/psychological-portrait.md</code>
the raw research	<code>prospects/{lead-id}/research-notes.md</code>
the witness verdict	<code>prospects/{lead-id}/witness-feedback.md</code>
to know why a lead was skipped	<code>results-summary.md</code> → pre-flagged; <code>prospects/{lead-id}/disqualify-check.md</code> → research-detected
to find leads in a persona/sector	<code>prospects/INDEX.md</code> — searchable table
a list of all campaigns	<code>campaigns/INDEX.md</code>
what notes were taken in a run	<code>craft/notes/{date}-*.md</code> — with the run's date
what patterns are now beliefs	<code>craft/patterns/</code>

I WANT...	WHERE IT IS
pending character changes	journal/proposed-character-diffs/
history of character changes	journal/applied-character-diffs/

12

output.json schema

The contract between the run and the human reviewer.

```
{
  "campaign_id": "string",
  "generated_at": "ISO-8601",
  "agent": "vendy",
  "results": [
    {
      "lead_id": "string",
      "status": "drafted | skipped | flag",
      "skip_reason": "string | null",
      "research_summary": {
        "linkedin_signal_strength": "high | med | low | none",
        "trigger_event": "string",
        "fallback_used": "individual | company | industry | none"
      },
      "draft": {
        "subject_line": "string",
        "body": "string",
        "voice_register_used": "string",
        "length_sentences": number
      },
      "witness": {
        "verdict": "ship | rewrite | flag",
        "what_works": "string",
        "what_fails": "string"
      },
      "ready_to_send": boolean,
      "human_review_required": boolean,
      "review_flags": ["string"]
    }
  ],
  "summary": {
    "total_leads": number,
    "drafted": number,
    "skipped": number,
    "flagged": number,
    "notes_added": number,
    "patterns_promoted": number
  }
}
```

The schema is also documented at `.claude/skills/run-campaign/references/output-schema.md`. Treat that file as authoritative if this PDF and the file diverge.

13

Five hard rules

Not preferences. The floor.

01 Do not edit self/.

Changes to the constitution flow through proposed diffs, human approval, and `/apply-diff`. There is no other path. Enforced by `pre-write-self-protect.py`.

02 Do not skip the witness.

Every draft is reviewed before it ships. The doer–witness separation is the load-bearing discipline of this architecture.

03 Do not ship a draft containing hard-blocked AI tells.

The deterministic detector is the floor; never write to bypass it. The list lives at `craft/cold-email/ai-tells-graveyard.md`.

04 Do not draft for a prospect you have not researched.

The portrait is the foundation; without it you are projecting, not seeing.

05 Do not chase replies dishonestly.

The work is the work. Wanting the reply is fine; corrupting the work to get it is not.

14

Phase roadmap

PHASE	WHEN	WHAT
Phase 1	now	Research → portrait → draft → witness review. Output JSON for human review and send. The current state.
Phase 2	next	Memory across campaigns. Promoted patterns inform future drafts. Cross-campaign retrieval. The library starts to compound.
Phase 3	horizon	Full autonomy. ICP definition → candidate sourcing → outreach → reply handling → follow-up cadence with guardrails.

"The work is the work. The next prospect deserves the same care."

— self/character.md

15

Privacy and data handling

What is gitignored

`experience/prospects/` contains personal data about real people. The folder is gitignored at the project root. Do not commit. Do not back up through standard project channels. If long-term retention is needed, use a separate, access-controlled mechanism with an explicit retention policy.

What is git-tracked

<code>self/</code>	The constitution and voice palette.
<code>craft/</code>	The accumulating library — no PII.
<code>experience/journal/</code>	Vendy's introspective record — about her, not them.
<code>experience/campaigns/</code>	Campaign records minus prospect-identifying inputs. README files in <code>prospects/</code> are tracked; the actual prospect folders are not.
<code>.claude/</code>	Agent definitions, skills, hooks, settings, lib.
<code>CLAUDE.md</code> , <code>README.md</code>	Project instructions and overview.

Disqualify signals

Vendy refuses to draft for a prospect in active crisis. The pre-check fires on the input's `disqualify_signals_check` array; the researcher also detects soft signals during work and writes `disqualify-check.md`. Categories include recent layoff, active breach disclosure, PTO indicator, public grief.

16

Setup and running it

What another machine needs to run Vendy. What runs without the optional pieces.

Prerequisites — required

Claude Code	CLI v2.0 or later. Vendy runs as a Claude Code project; this is the runtime.
The Vendy repo	Cloned to a working directory. <code>git clone</code> then <code>cd</code> into it before launching Claude Code.
Python 3.9+	The six hooks under <code>.claude/hooks/</code> include Python scripts that run via the OS interpreter.

Prerequisites — optional

Chrome / Arc / Brave	A modern Chromium-based browser. Needed only if you want authenticated LinkedIn deep-reads.
claude-in-chrome-mcp	MCP server binary referenced in <code>.mcp.json</code> . When installed on PATH and approved per-user, the researcher drives your logged-in Chrome for LinkedIn, signed-in GitHub, and other authenticated reads.

Without the optional MCP — fallback mode

Vendy runs end-to-end without `claude-in-chrome-mcp`. The researcher sub-agent falls back to `WebFetch` + `WebSearch` for what is publicly fetchable, plus aggregator sites (ZoomInfo, ContactOut, Wiza) for career-history reconstruction. The depth is lower — no comments, no reactions, no posts beyond what LinkedIn serves logged-out — but the workflow is intact.

16

Setup — running a session

Approving the MCP server

The repo's `.mcp.json` declares the server. Each new user must approve it themselves — approval lands in `~/.claude/settings.local.json`, which is gitignored.

```
$ claude mcp list
  claude-in-chrome: claude-in-chrome-mcp  ✓ Connected

# If not connected, accept the prompt on first session start, or run:
$ claude mcp approve claude-in-chrome
```

Starting a session

```
$ cd path/to/vendy
$ claude

# The SessionStart hook prints the welcome banner and injects:
#   - self/character.md
#   - the latest experience/journal/ entry
#   - craft/trends/current.md
```

Before the first real campaign

<code>/refresh-trends</code>	Pulls current security/eng narratives into <code>craft/trends/current.md</code> . Auto-runs at campaign start if older than 14 days; manual run advisable on a fresh checkout.
<code>craft/exemplars/</code>	Optionally seed with 5–10 of your team's best past cold emails using the schema in <code>exemplars/format.md</code> . Vendy reads exemplars during drafting.

Running a campaign

```
# 1. Write input JSON. Schema:
#   .claude/skills/run-campaign/references/input-schema.md

# 2. Invoke: /run-campaign path/to/your-input.json

# 3. Read outputs at:
#   experience/campaigns/{campaign-id}/output.json
#   experience/campaigns/{campaign-id}/results-summary.md
#   experience/prospects/{lead-id}/final.md (per lead)
```

16

Setup — verifying and current state

Quick verification checklist

<code>which claude</code>	Path to the claude binary; version 2.0 or later.
<code>python3 --version</code>	3.9 or later.
<code>claude mcp list</code>	<code>claude-in-chrome</code> ✓ <code>Connected</code> if the optional MCP is installed; otherwise the line is missing or shows an error.
Open the project	Session banner appears; <code>self/character.md</code> visible in context.
<code>/refresh-trends</code>	Writes a fresh <code>craft/trends/current.md</code> with today's date in frontmatter.

Current state on this machine

A snapshot taken at the time of generation, for the operator's reference.

Claude Code	<code>2.1.128</code> at <code>~/.local/bin/claude</code> ✓
Python	<code>3.9.6</code> (system) ✓
Chrome / Arc	Both installed in <code>/Applications/</code> ✓
<code>claude-in-chrome-mcp</code>	Not installed. Researcher operates in fallback mode (WebFetch + WebSearch + aggregators).
Real research depth	See <code>experience/prospects/justin-stanford-coalesce/research-notes.md</code> — career history reconstructed from ZoomInfo, ContactOut, Wiza cross-confirmed.

PRIVACY `experience/prospects/` contains personal data and is gitignored. Do not commit it; do not back it up through standard project channels. See section 15.

17

Colophon

Vendy runs as a Claude Code project on a developer's machine. The character is one document. The library is files on disk. The discipline is the doer–witness separation. The rest is just good writing.

Stack

Runtime	Claude Code · Claude Opus 4.6 / 4.7
Browser automation	Claude in Chrome MCP server
Memory	Filesystem — self/, craft/, experience/
Privacy	experience/prospects/ is gitignored
Configuration	.claude/settings.json
Local overrides	.claude/settings.local.json (gitignored)
MCP config	.mcp.json

How to extend

Add a persona	Write <code>craft/personas/{role}.md</code> in the format from <code>personas/format.md</code> .
Add an exemplar	Drop a dissected past email into <code>craft/exemplars/</code> per the format file.
Add a research playbook	Add a markdown file under <code>craft/research-methodology/</code> ; the researcher will load it when relevant.
Change Vendy's character	Write a proposed diff under <code>experience/journal/proposed-character-diffs/</code> ; add approval frontmatter; run <code>/apply-diff</code> .

